# Predicting NFL Margins with Gaussian Processes: A Replication and Extension of Warner (2010)

Nate Krall & Viraj Acharya

2026-04-26

## Table of contents

**1 Background** **2**

**2 Research Question and Previous Results** **2**

**3 Data** **3**
3.1 Source and Scope . . . 3
3.2 Features . . . 3
3.3 Train / Test Split . . . 4

**4 Methods** **4**
4.1 Gaussian Process Regression . . . 4
4.2 Kernel Choice . . . 5
4.3 Computed Strength Rankings . . . 5
4.4 Feature Selection . . . 6
4.5 Betting Framework . . . 6
4.6 Random Forest Comparison . . . 7

**5 Results** **8**
5.1 Era Comparison: Has Vegas Gotten Sharper? . . . 8
5.2 Feature Selection Results . . . 9
5.3 Betting Framework . . . 9
5.4 Random Forest Betting Results . . . 10

**6 Discussion** **11**
6.1 Key Takeaways . . . 11
6.2 Limitations . . . 12

**7 Conclusion** **12**

**8 References** **12**

**9 Appendix: Code** **13**
9.1 Data Loading and Feature Construction . . . 13

9.2 Computed Strength (Keener Ranking) . . . . . 13
9.3 GP Model . . . . . 14
9.4 Betting Framework . . . . . 14

# 1 Background

Jim Warner in *"Predicting Margin of Victory in NFL Games: Machine Learning vs. the Las Vegas Line"* (2010) asks whether a statistical model trained purely on box-score team statistics can predict NFL game margins of victory, and whether it can compete with Vegas point spreads. Sportsbook spread lines are essentially a prediction on the favored team's margin of victory: they propose an amount of points they believe the favored team is projected to win by, and they allow bettors to place money on the favored team to win by more points than that proposition or by less points (including a loss altogether). By creating a model that outputs a prediction on the home team's margin of victory, we naturally have a metric to compare to Vegas's spread proposition in order to trigger a bet. What makes this modeling problem particularly challenging is that a win rate of greater than 50% over the Vegas line is not necessarily sufficient; since sportsbooks charge a small fee embedded into the odds of every bet attempting to create an edge, our model must pick the correct side of the Vegas line around 52.4% or greater to ensure profitability, assuming the standard vig (represented by -110 in American odds) that is widespread across sportsbooks.

Warner's paper, and our project thereafter, analyze the existence of a model capable of achieving this win rate. This question is important because, if it is true that a machine learning model trained on accessible data can consistently outperform the Las Vegas Line over 52.4% of the time, then we would have a profitable betting scheme we could use to exploit sportsbooks.

---

# 2 Research Question and Previous Results

Using a Gaussian Process (GP) regressor on data from the 2000–2009 NFL seasons, Warner found that his model achieved roughly 64% winner accuracy compared to Vegas's 68.5%, with comparable mean absolute error on margin of victory (~11.5 vs. ~11.3 points). Warner was unable to produce a win rate over 52.4% versus the Vegas line, and concluded that a GP trained on box-score data is insufficient in achieving profitability on sportsbooks. Since it has been quite some time since the paper's release, we aim to apply Warner's methods on modern NFL data to see how the conclusions have changed. The dynamics of the NFL are far different today than they were well over a decade ago, and sportsbooks have become far more robust in recent years, so we anticipate our conclusions will differ from those of the paper in some fashion. We replicate Warner's methodology on the years 2015–2024 (a 9 year span like the paper used) to ask a new question: has Vegas gotten sharper in the 15 years since the paper was published? If the GP's accuracy has stayed flat while Vegas has improved, that would suggest the sportsbook market has become more informationally efficient over time.

---

# 3 Data

## 3.1 Source and Scope

Our analysis uses two publicly available datasets from the `nflverse` project, an open-source repository of NFL data. Game-level results and Vegas point spreads come from `nflverse`'s `games.csv`, which contains every NFL regular-season game since 1999 along with final scores, point spreads, and totals. Per-game team statistics — yards, turnovers, and points — are aggregated from the corresponding season-level play-by-play files hosted on the `nflverse-data` GitHub releases. Our analysis covers the 2015–2024 regular seasons, yielding 2,623 total games. After filtering out early-season games that lack sufficient prior-game history to compute rolling features, 1,885 matchups remain in the modeling dataset.

The `spread_line` column in `games.csv` is sourced from ESPN and represents the consensus spread at or near game time. Warner (2010) does not specify whether his spread data reflects the opening line, the closing line, or some consensus thereof, and `nflverse` does not explicitly label its line as a closing line either. This is a minor limitation worth noting: closing lines tend to be sharper than opening lines as they incorporate late-breaking information such as injury reports and sharp betting activity, so it is possible our Vegas benchmark and Warner's are not perfectly comparable on this dimension. However, all lines from open to close tend to be very similar if not the same, so our results still hold significance even if our lines and the paper's lines were snapshotted at different times.

## 3.2 Features

We construct 46 features per matchup (23 per team) from the play-by-play data, mirroring and extending Warner's original feature set. Warner used 10 base statistics per team, which are points scored, points allowed, total yards, total yards allowed, rushing yards, rushing yards allowed, passing yards, passing yards allowed, turnovers taken, and turnovers lost. For each of these 10 statistics we also compute two versions: a season-to-date average, which captures a team's overall performance level for the year, and a 4-game rolling average (referred to as a streak), which captures recent form and momentum. We also compute season-to-date and 4-game rolling win percentage for each team, and the Keener eigenvalue-based computed strength rating described in Section 4.4. This yields 23 features per team and 46 total per matchup.

All features are computed using only games played prior to the current week to prevent data leakage, meaning the model only ever sees information that would have been available at kickoff. The 4-game rolling features require a team to have played at least 4 prior games in the current season, which means matchups from weeks 1–4 of each season are excluded from the dataset entirely. This is the reason our 2,623 loaded games reduce to 1,885 usable matchups. The table below summarizes how the 46 features per matchup are dispersed.

| Feature Group | Variables | Per Team |
|---|---|---|
| Season-to-date averages | Points scored/allowed, total/rush/pass yards, turnovers taken/lost | 10 |
| 4-game rolling averages | Same 10 stats as above | 10 |

| Feature Group | Variables | Per Team |
|---|---|---|
| Win percentage | Season-to-date and 4-game rolling | 2 |
| Computed strength | Keener eigenvalue ranking (Section 4.4) | 1 |
| **Total** | | $\mathbf{23 \times 2 = 46}$ |

### 3.3 Train / Test Split

We train our Gaussian Process model on this data from the years 2015 to 2022, and we test the model on games from 2023 and 2024. This mirrors the scheme Warner employed in the paper and allows us to achieve directly comparable results while maintaining a sufficient amount of data for both training and testing purposes.

---

## 4 Methods

### 4.1 Gaussian Process Regression

A Gaussian Process (GP) is a non-parametric Bayesian regression model that represents a probability distribution over functions rather than a single fitted function. Instead of committing to a fixed functional form, a GP considers every function that could plausibly explain the training data and weights them by how well they fit. The core modeling assumption is simple and intuitive, namely that similar inputs should produce similar outputs, which in our setting means that two matchups with similar team statistics should produce similar predicted margins of victory. The mechanism that defines what "similar" means is the kernel, discussed in Section 4.2.

The most important practical difference between a GP and linear regression is what each returns for a prediction. Linear regression produces a point estimate, outputting a single number for the predicted margin and treating uncertainty as a global, constant quantity inherited from the residual variance of the fit. A GP instead returns a full predictive distribution at every test point, outputting a posterior mean for the predicted margin alongside a posterior standard deviation that captures a calibrated, matchup-specific uncertainty estimate. The uncertainty grows in regions of feature space where the training data is sparse and shrinks where many similar matchups have been observed. This locally-aware uncertainty is something a linear model cannot natively provide, and it is the property that makes the GP particularly well-suited to Warner's betting framework, which uses the model's confidence to decide whether a disagreement with Vegas is worth acting on.

The flexibility of a GP comes at a cost, as fitting a GP requires inverting an $N \times N$ kernel matrix. This gives training a computational complexity of $O(N^3)$, whereas linear regression scales as $O(N)$. This makes GPs best suited to small and medium sized datasets where precision and uncertainty quantification matter more than raw speed, a profile that fits our roughly 1,800-game training set well. For applications with millions of training points, or where interpretability of individual coefficients is paramount, a parametric model is the better choice. However, in our case for predicting

NFL margins from a modest pool of team statistics, where every prediction also needs an honest sense of how confident the model actually is, a GP is the natural fit.

### 4.2 Kernel Choice

The kernel is the function that effectively computes a similarity score between any two input points, which determines how much their outputs should be correlated under the prior. The choice of kernel is where the modeling assumptions live, encoding beliefs about the smoothness and structural properties of the underlying function being learned. Different kernels can also be added together to combine assumptions about signal and noise, which is exactly what we do here.

We use the same kernel as Warner, a squared-exponential (RBF) kernel plus a WhiteKernel to capture irreducible noise.

$$k(x_i, x_j) = \sigma^2 \exp\left(-\frac{\|x_i - x_j\|^2}{2\ell^2}\right) + \sigma_n^2 \delta_{ij}$$

The RBF term assigns a similarity value between two matchups based on the squared distance between their feature vectors, returning a value close to 1 when inputs are nearby and decaying toward 0 as they grow further apart. The length scale $\ell$ controls how quickly that decay happens, where a small $\ell$ means even nearby points in feature space are treated as essentially independent and a large $\ell$ means the model expects the underlying function to vary smoothly across inputs. The WhiteKernel adds a constant noise variance $\sigma_n^2$ along the diagonal of the kernel matrix, encoding the assumption that even two identical inputs would not produce identical outputs because of irreducible randomness from injuries, weather, individual plays, and other sources of variance in football outcomes.

Both $\ell$ and $\sigma_n^2$ are learned by maximizing the log marginal likelihood on the training data, and the fitted kernel for the Warner 4-feature GP comes out to RBF($\ell = 7.31$) + WhiteKernel($\sigma_n^2 = 0.872$). Since features are standardized to unit variance before fitting, the length scale of 7.31 is in standardized feature units and indicates a smooth fit where correlation between matchups drops off only when they differ substantially across multiple features at once. The noise level of 0.872 is in standardized response units and represents a meaningful share of the total variance in margins of victory, which is exactly what we should expect since box-score statistics cannot account for all possible sources of variance that can meaningfully affect the outcome of a game.

### 4.3 Computed Strength Rankings

Following Warner Section 2.2.1, we compute an eigenvalue-based strength rating for each team at each week of the season using Keener's method. The motivation is that a simple win-loss record fails to account for opponent quality, since beating a strong team should weigh more heavily than beating a weak one, and a raw record cannot distinguish between the two. Keener's method addresses this by defining each team's strength recursively as a weighted average of the strengths of its opponents, with weights determined by how it fared in each matchup. Concretely, we build a matrix $A$ where each entry $A_{ij}$ encodes how decisively team $i$ has outperformed team $j$ in their head-to-head history, smoothed by the function $h(x) = 0.5 + 0.5 \cdot \text{sgn}(x - 0.5) \cdot \sqrt{|2x - 1|}$. This dampens blowout wins, rewarding the act of winning a matchup more heavily than running up the

score and preventing a team's rating from being inflated by lopsided victories over weak opponents. The recursive definition gives rise to an eigenvalue equation, and each team's strength rating is its corresponding entry in the dominant eigenvector of $A$, computed efficiently via power iteration.

### 4.4 Feature Selection

With 46 available features, using all of them simultaneously is not necessarily optimal. More features increases the dimensionality of the input space, which can hurt GP performance as the kernel's notion of similarity between two matchups becomes diluted across many variables, some of which may be redundant or noisy. Warner addressed this by running a forward feature selection procedure on his training data to identify a compact, high-signal feature subset. Warner's selection process began with a base set of home and away win percentage, treating these as the two most fundamental predictors of team quality, which is a natural conclusion. From there, he evaluated every remaining candidate feature by adding it to the base set one at a time and measuring the resulting model's cross-validated mean absolute error (MAE). He retained the top candidates from this screening step and then ran a greedy forward search: iteratively adding whichever feature produced the largest reduction in cross-validated MAE, stopping when no remaining feature improved performance. Warner's search converged on a final set of four features: home win percentage, away win percentage, away computed strength, and away passing yards per game (4-game streak). Notably, three of the four are away-team statistics.

We run the GP using Warner's exact four features on our 2015–2024 data as our primary replication, since this allows us to compare our results directly to the numbers reported in the paper. However, we also re-run the feature selection procedure on our own training data to find the optimal feature subset for the modern era. This serves two purposes: it tells us whether Warner's 4-feature set still represents a good choice on modern data, and it allows our modern GP to make predictions based on features calibrated for the modern game. Our selection process mirrors Warner's exactly. We start from the same base set of home and away win percentage, screen all 44 remaining candidates using leave-one-season-out cross-validation on the 2015–2022 training seasons, take the top 20 by CV MAE as the search pool, and run a greedy forward search over those 20, stopping when no additional feature reduces cross-validated error.

Our search converged on eight features:

- Home win percentage, away win percentage (base set)
- Away points scored (season-to-date average)
- Home computed strength
- Away total yards (season-to-date average and 4-game streak)
- Home points allowed (4-game streak)
- Away turnovers lost (season-to-date average)

### 4.5 Betting Framework

One of the key advantages of a Gaussian Process over a simpler regression model is that every prediction comes with a calibrated uncertainty estimate — a standard deviation $\sigma$ reflecting how confident the model is about that particular matchup. Warner's betting framework exploits this directly. Rather than betting on every game, Warner proposes placing a bet only when the Vegas

spread line falls outside the GP's 95% confidence interval around its predicted margin. The intuition is that a disagreement between the model and Vegas that falls within the model's own uncertainty is not worth acting on, while a disagreement large enough to exceed the 95% interval represents a more genuine signal that the line may be mispriced.

Concretely, let $\hat{y}$ be the GP's predicted home margin, $\sigma$ be its predicted standard deviation, and $s$ be the Vegas spread line. Warner bets on the home team to cover when $\hat{y} - 1.96\sigma > s$, and on the away team when $\hat{y} + 1.96\sigma < s$. When the spread line falls inside the 95% interval, no bet is placed. This is illustrated in the figure below.

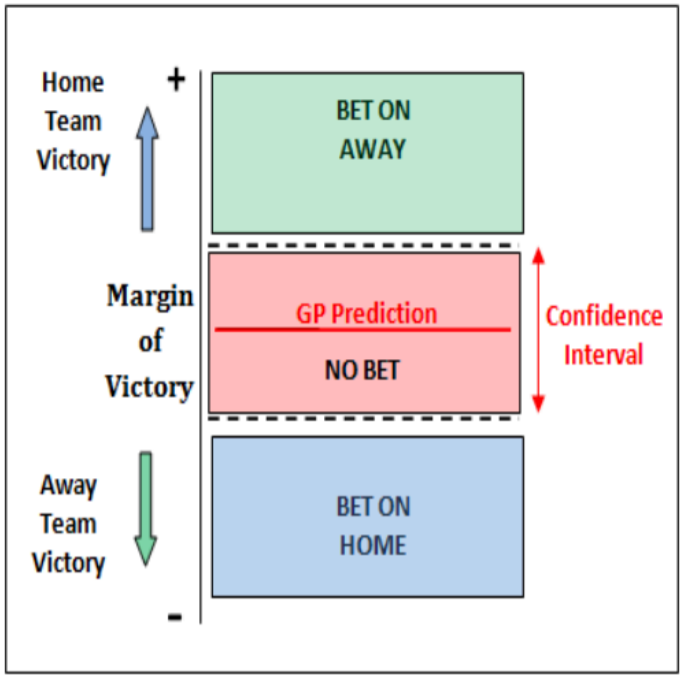


Figure 1: Betting strategy diagram: a bet is placed only when the Vegas spread line falls outside the GP's confidence interval around its prediction.

As an extension to Warner's framework, we generalize the fixed 95% threshold to a tunable parameter $k$, evaluating the strategy across $k$ values from 0.25 to 1.5. A lower $k$ corresponds to a narrower confidence interval and triggers more bets, including cases where the model's confidence is modest; a higher $k$ filters down to only the most emphatic disagreements between the model and Vegas but leaves very few bets in the sample. This allows us to examine the tradeoff between bet volume and selectivity, and to assess whether any threshold produces a profitable win rate. As a reminder, at standard $-110$ vig, the break-even win rate required for profitability is 52.4%. We apply this framework to both the Warner 4-feature GP and our 8-feature GP and report the results in Section 5.3.

### 4.6 Random Forest Comparison

A natural question when evaluating the GP is whether its performance is a property of the model class or simply a reflection of the underlying data. If the features we constructed are genuinely not predictive enough to beat Vegas, then any reasonable regression model should struggle similarly.

To investigate this, we also fit a Random Forest regressor on the same 8 selected features and the same 2015–2022 training data.

Random Forests are an ensemble method that averages predictions across a large number of decision trees, each trained on a bootstrapped sample of the data. Unlike a GP, a Random Forest does not assume any particular functional form for the relationship between features and outcomes, making it a flexible alternative that serves as a useful point of comparison. We trained 500 trees with a minimum leaf size of 5 to prevent overfitting, and we estimated prediction uncertainty as the standard deviation of predictions across the 500 trees — a natural analogue to the GP's posterior standard deviation. This allows us to apply Warner's betting framework to the Random Forest using the exact same rules, and compare the betting results of the two models directly.

---

# 5 Results

## 5.1 Era Comparison: Has Vegas Gotten Sharper?

We begin by evaluating the GP using Warner's exact four features from the paper, which are home win percentage, away win percentage, away computed strength, and away passing yards per game (4-game streak), on our 2015–2024 data. This allows us to make a direct apples-to-apples comparison with the numbers reported in the paper.

On the 2023–2024 test set, our GP achieved a mean absolute error of 9.98 points and a winner accuracy of 64.4%. The MAE is a notable improvement over Warner's reported ~11.5 points, but this does not necessarily mean the model has gotten better at predicting NFL outcomes. A more likely explanation is that the league itself has become more competitive over time, with games being decided by smaller margins on average, making the raw prediction error smaller regardless of model quality. The winner's accuracy tells a cleaner story: at 64.4%, it is identical to what Warner reported in 2010, suggesting the model's ability to pick the correct side has not meaningfully changed in 15 years.

Vegas, however, has improved substantially. The spread line achieved a winner accuracy of 72.1% on our test set, up from ~68.5% in Warner's data. The gap between the GP and Vegas has widened from roughly 4 percentage points to nearly 8. This widening suggests that the GP has stood still, the sportsbook market has become significantly more informationally efficient. The era comparison is summarized in the figure below.

**Era Comparison: GP vs. Vegas (2008–09 → 2023–24)**

| | GP (Warner 2008–09) | Vegas (Warner 2008–09) | GP (Ours 2023–24) | Vegas (Ours 2023–24) |
|---|---|---|---|---|
| MAE (margin of victory) | ~11.5 | ~11.3 | 9.98 | **9.39** |
| Winner accuracy | 64.4% | **68.5%** | 64.4% | **72.1%** |

Figure 2: Era comparison: GP vs. Vegas across two data windows.

## 5.2 Feature Selection Results

It is not fair to conclude GP performance has not changed by using the exact same feature set as the paper, since certain features are plausibly more or less predictive in the modern football era. To address this, we created a new GP trained on the 8-feature set discussed above that is supposedly better calibrated to the modern game. Running the GP on this updated feature set improved winner accuracy from 64.4% to 67.1% on the 2023–2024 test set, a meaningful gain that suggests our selected features do carry more signal for the current era. However, MAE was essentially unchanged at 10.02 points versus 9.98 with Warner's features, indicating that the additional features help the model identify the correct side of a matchup more often, but do not help it predict the margin of victory more precisely. Even with the improved feature set, the GP still falls well short of Vegas's 72.1% winner accuracy, reinforcing that the gap is not primarily a feature selection problem. The results are summarized in the figure below.

**Feature Selection: Warner 4 Features vs. Our 8-Feature Set**

| | Warner 4-Feature GP (2023–24 test) | Our 8-Feature GP (2023–24 test) | Vegas (2023–24 test) |
|---|---|---|---|
| MAE (margin of victory) | 9.98 | 10.02 | **9.39** |
| Winner accuracy | 64.4% | **67.1%** | **72.1%** |

Figure 3: Feature selection: Warner's 4 features vs. our 8-feature set.

## 5.3 Betting Framework

We sweep $k$ from 0.25 to 1.5 to evaluate whether any confidence threshold produces a win rate above the 52.4% break-even win rate. Unsurprisingly, neither the Warner 4-feature GP nor our 8-feature GP achieves a profitable win rate at any threshold with a meaningful number of bets (more than just a few). What is surprising is that no bets were placed for $k > 1.25$, while Warner's GP placed several bets with $k = 1.96$. Looking at the 4-feature GP for example, the median

posterior standard deviation is approximately 13.24 points, so a bet at $k = 1.96$ would only be triggered when the spread line differs from the GP's prediction by more than 26 points, which is an extraordinary disagreement that almost never occurs in practice. So, why did this disagreement occur so often back in 2008-9? Precisely because Vegas was less accurate at the time; a less efficient market produced more lines that appeared exploitable relative to the model's uncertainty. Today, Vegas prices games tightly enough that the GP almost never sees a spread it disagrees with by that margin, leaving very little room for the betting framework to operate. The betting results for both models are shown in the figure below.

**Betting Framework: k-Threshold Results (break-even = 52.4%)**

**Warner 4-Feature GP (2023–24)**

| | Total Bets | Wins | Win Rate | Profitable? |
|---|---|---|---|---|
| k = 0.25 | 158 | 75 | 47.5% | no |
| k = 0.50 | 39 | 20 | 51.3% | no |
| k = 0.75 | 5 | 4 | 80.0% | YES |
| k = 1.00 | 1 | 0 | 0.0% | no |
| k = 1.25 | 0 | 0 | — | no |
| k = 1.50 | 0 | 0 | — | no |

**Our 8-Feature GP (2023–24)**

| | Total Bets | Wins | Win Rate | Profitable? |
|---|---|---|---|---|
| k = 0.25 | 159 | 77 | 48.4% | no |
| k = 0.50 | 37 | 18 | 48.6% | no |
| k = 0.75 | 12 | 7 | 58.3% | YES |
| k = 1.00 | 1 | 1 | 100.0% | YES |
| k = 1.25 | 0 | 0 | — | no |
| k = 1.50 | 0 | 0 | — | no |

*Standard −110 vig · Home covers when result > spread_line · Away covers when result < spread_line*

Figure 4: Betting framework: k-threshold win rates for both models.

## 5.4 Random Forest Betting Results

Random Forests are widely regarded as a strong baseline for tabular data, so applying the same betting framework to the RF provides a meaningful check on whether the GP's failure to find a profitable edge is a model-specific limitation or a property of the data itself. One notable difference between the two models is that the RF places more bets at every $k$ level. This is because the RF's uncertainty estimate — the standard deviation across 500 trees — tends to be smaller and more tightly calibrated than the GP's posterior $\sigma$, so the spread line more frequently falls outside the RF's confidence interval and triggers a bet. Despite placing more bets, the RF does not achieve a consistent profitable win rate at any threshold either. The fact that two models with very different architectures, assumptions, and uncertainty estimates both fail to find a reliable edge against the spread reinforces the conclusion that the problem lies with the data rather than the model: box-score statistics simply do not contain enough information to systematically exploit modern Vegas lines.

**GP vs. Random Forest: Betting Results (break-even = 52.4%)**

**8-Feature GP (2023–24)**

| | Total Bets | Wins | Win Rate | Profitable? |
|---|---|---|---|---|
| k = 0.25 | 159 | 77 | 48.4% | no |
| k = 0.50 | 37 | 18 | 48.6% | no |
| k = 0.75 | 12 | 7 | 58.3% | YES |
| k = 1.00 | 1 | 1 | 100.0% | YES |
| k = 1.25 | 0 | 0 | — | no |
| k = 1.50 | 0 | 0 | — | no |

**Random Forest (same 8 features)**

| | Total Bets | Wins | Win Rate | Profitable? |
|---|---|---|---|---|
| k = 0.25 | 264 | 135 | 51.1% | no |
| k = 0.50 | 135 | 65 | 48.1% | no |
| k = 0.75 | 49 | 25 | 51.0% | no |
| k = 1.00 | 20 | 10 | 50.0% | no |
| k = 1.25 | 6 | 3 | 50.0% | no |
| k = 1.50 | 3 | 2 | 66.7% | YES |

*RF uncertainty = std across 500 trees · Same betting rules as GP framework*

Figure 5: GP vs. Random Forest betting results across k thresholds.

---

# 6 Discussion

## 6.1 Key Takeaways

The most pronounced finding from our work is that Vegas has gotten substantially sharper in the 15 years since Warner's paper was published. Winner accuracy on the spread line jumped from roughly 68.5% in 2008–09 to 72.1% in 2023–24, while our GP's winner accuracy stayed flat at exactly the 64.4% Warner reported. The gap between the model and Vegas widened from about 4 percentage points to nearly 8, which suggests that the sportsbook market has become meaningfully more informationally efficient over time even as a GP trained on box-score statistics has not improved alongside it.

The second finding is that neither letting the model choose its own features nor swapping in a different model class meaningfully closes that gap. Our 8-feature GP, selected via the same forward search procedure but using modern training data, picks winners at 67.1% compared to Warner's 4-feature set at 64.4%, but MAE remains essentially unchanged at around 10 points. The additional features help the model identify the correct side of a matchup more often, but they do not improve its ability to predict the actual margin of victory. The random forest, despite making very different assumptions about the structure of the data and producing a different uncertainty profile, lands in a similar place. This consistency across model classes points to a feature problem rather than a model problem, since two architectures with different inductive biases arrive at roughly the same ceiling on the same data.

The third finding follows naturally from the first two. The betting framework does not produce a profitable edge at any threshold with a meaningful sample size, with win rates falling below the 52.4% break-even line whenever enough bets are placed to draw a reliable conclusion. Beyond a small confidence threshold the model essentially never disagrees with Vegas by enough to justify a

bet, since the sportsbook prices games tightly enough that the kinds of large mispricings Warner could occasionally exploit in 2008–09 simply no longer exist in our test window.

### 6.2 Limitations

Several limitations of our analysis are worth noting. The most important is that our feature set is restricted to box-score team statistics and contains no situational or contextual information that real sportsbook models almost certainly use. We do not include injury reports, which can dramatically alter the expected outcome of a game when a starting quarterback or another key player is ruled out late in the week. We do not include weather conditions, which can swing scoring, totals, and margins in outdoor games played in cold or windy environments. We also do not include roster turnover information from offseason trades, free agency, or coaching changes, all of which can render a team's prior-season statistics misleading. Another notable limitation concerns the spread line itself, since the data from nflverse is sourced from ESPN at or near game time and we treat it as a single static value, while in practice spreads move throughout the week as injuries are disclosed and as sharp money enters the market. Incorporating line movement, the divergence between opening and closing lines, or differences across sportsbooks could provide useful signal that our model does not currently see. Vegas lines are very likely informed by all of these factors, which gives the sportsbook an information advantage that no purely box-score model can match.

---

## 7 Conclusion

Warner concluded in 2010 that a Gaussian Process trained on box-score statistics could not consistently beat the Las Vegas line, and our replication on 2015–2024 data reaches the same conclusion in stronger form. Vegas has gotten meaningfully sharper over the intervening 15 years while models driven by box-scores have not, leaving an even smaller window for the kind of disagreements Warner's betting framework was designed to exploit. Whatever edge once existed for a model of this kind has effectively closed, and beating modern sportsbook lines almost certainly requires the situational and contextual information that team statistics alone cannot capture.

---

## 8 References

Warner, G. (2010). *Winning isn't Everything: A Contextual Analysis of Results in the NFL Using Gaussian Processes.* Proceedings of the MIT Sloan Sports Analytics Conference.

---

# 9 Appendix: Code

The complete analysis code is available in `nfl_gp_starter.ipynb`. Key sections are reproduced below for reference.

## 9.1 Data Loading and Feature Construction

```
import pandas as pd
import numpy as np

TRAIN_SEASONS = list(range(2015, 2023))
TEST_SEASONS  = list(range(2023, 2025))
ALL_SEASONS   = TRAIN_SEASONS + TEST_SEASONS

BASE_STATS = [
    "points_scored", "points_allowed",
    "total_yards", "total_yards_allowed",
    "rush_yards", "rush_yards_allowed",
    "pass_yards", "pass_yards_allowed",
    "turnovers_taken", "turnovers_lost",
]

games = pd.read_csv(
    "https://raw.githubusercontent.com/nflverse/nfldata/master/data/games.csv",
    low_memory=False
)
games = games[
    (games["game_type"] == "REG") &
    (games["result"].notna()) &
    (games["season"].isin(ALL_SEASONS))
].copy()
```

## 9.2 Computed Strength (Keener Ranking)

```
def h_func(x):
    return 0.5 + 0.5 * np.sign(x - 0.5) * np.sqrt(np.abs(2 * x - 1))

def compute_strength(games_df):
    teams = sorted(set(games_df["home_team"]) | set(games_df["away_team"]))
    n = len(teams)
    idx = {t: i for i, t in enumerate(teams)}
    S = np.zeros((n, n))
    for _, row in games_df.iterrows():
```

```
        i, j = idx[row["home_team"]], idx[row["away_team"]]
        S[i][j] += row["home_score"]
        S[j][i] += row["away_score"]
    A = np.zeros((n, n))
    for i in range(n):
        for j in range(n):
            total = S[i][j] + S[j][i]
            if total > 0:
                A[i][j] = h_func((S[i][j] + 1) / (total + 2))
    r = np.ones(n) / n
    for _ in range(200):
        r_new = A @ r
        norm = r_new.sum()
        if norm == 0: break
        r_new /= norm
        if np.allclose(r, r_new, atol=1e-10): break
        r = r_new
    return dict(zip(teams, r))
```

### 9.3 GP Model

```
from sklearn.gaussian_process import GaussianProcessRegressor
from sklearn.gaussian_process.kernels import RBF, WhiteKernel
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import mean_absolute_error

GP_FEATURES = ["H_win_pct", "A_win_pct", "A_computed_strength", "A_pass_yards_streak"]

scaler  = StandardScaler()
X_train = scaler.fit_transform(train_gp[GP_FEATURES])
X_test  = scaler.transform(test_gp[GP_FEATURES])

kernel = RBF() + WhiteKernel()
gp = GaussianProcessRegressor(kernel=kernel, n_restarts_optimizer=5, normalize_y=True)
gp.fit(X_train, train_gp["result"].values)

y_pred, y_std = gp.predict(X_test, return_std=True)
```

### 9.4 Betting Framework

```
def evaluate_betting(test_df, k):
    home_bets = test_df[test_df["gp_pred"] - k * test_df["gp_std"] > test_df["spread_line"]]
    away_bets = test_df[test_df["gp_pred"] + k * test_df["gp_std"] < test_df["spread_line"]]
```

```
    home_wins = (home_bets["result"] > home_bets["spread_line"]).sum()
    away_wins = (away_bets["result"] < away_bets["spread_line"]).sum()
    total_bets = len(home_bets) + len(away_bets)
    total_wins = int(home_wins + away_wins)
    win_rate   = total_wins / total_bets if total_bets > 0 else float("nan")
    return {"k": k, "bets": total_bets, "wins": total_wins, "win_rate": win_rate}
```